

Formula

$$AX = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \dots + a_{1n}x_n \\ a_{21}x_1 + a_{22}x_2 + \dots + a_{2n}x_n \\ \vdots \\ a_{m1}x_1 + a_{m2}x_2 + \dots + a_{mn}x_n \end{bmatrix}$$

relevant elements

selected elements

How many selected items are relevant?

Precision =

How many relevant items are selected?

Recall =

sigmoid: $\frac{1}{1+e^{-x}}$
 sigmoid drv: $\sigma(1-\sigma)$
 tanh drv: $1 - \tanh^2(x)$

Batch Norm

Batch Normalization

$$\mu = \frac{1}{m} \sum_i x^{(i)}$$

$$\sigma^2 = \frac{1}{m} \sum_i (x^{(i)} - \mu)^2$$

$$z_{\text{norm}}^{(i)} = \frac{x^{(i)} - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$\hat{x}^{(i)} = \gamma z_{\text{norm}}^{(i)} + \beta$$

- Mini-batch Batch Norm
- Usually applied before activation
- At test time: use running average of mean and var computed during train time

(i) accelerates learning by reducing co-variate shift, decoupling dependence of layers, and/or allowing for higher learning rates/ deeper networks, (ii) accelerates learning by normalizing contours of output distribution to be more uniform across dimensions, (iii) Regularizes by using batch To receive full credit, these responses included three components:

(i) Mini-batches might be small at test time. (ii) Smaller mini-batches mean the mini-batch statistics are more likely to

differ at training. (iii) Moving averages are better estimates.

Optimizers

Algorithm 1: Adam, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-9}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize
Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates
Require: $f(\theta)$: Stochastic objective function with parameters θ
Require: θ_0 : Initial parameter vector
 $m_0 \leftarrow 0$ (Initialize 1st moment vector)
 $v_0 \leftarrow 0$ (Initialize 2nd moment vector)
 $t \leftarrow 0$ (Initialize timestep)
while θ_t not converged **do**
 $t \leftarrow t + 1$
 $g_t \leftarrow -\nabla_{\theta} f(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)
 $m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)
 $v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)
 $\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)
 $\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)
 $\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)
end while
return θ_t (Resulting parameters)

There are a few important differences between RMSProp with momentum and Adam: RMSProp with momentum generates its parameter updates using a momentum on the rescaled gradient, whereas Adam updates are directly estimated using a running average of first and second moment of the gradient. RMSProp also lacks a bias-correction term; this matters most in case of a value of 2 close to 1 (required in case of sparse gradients), since in that case not correcting the bias leads to very large stepsizes and often divergence,

Backprop

$$Z^{[1]} = \begin{bmatrix} z_1^{[1]} \\ z_1^{[2]} \end{bmatrix} = \begin{bmatrix} w_{11}^{[1]} & w_{12}^{[1]} \\ w_{21}^{[1]} & w_{22}^{[1]} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad A^{[1]} = \begin{bmatrix} a_1^{[1]} \\ a_2^{[1]} \end{bmatrix} = \begin{bmatrix} \sigma(z_1^{[1]}) \\ \sigma(z_1^{[2]}) \end{bmatrix}$$

$$Z^{[2]} = \begin{bmatrix} z_2^{[1]} \\ z_2^{[2]} \end{bmatrix} = \begin{bmatrix} w_{11}^{[2]} & w_{12}^{[2]} \\ w_{21}^{[2]} & w_{22}^{[2]} \end{bmatrix} \begin{bmatrix} a_1^{[1]} \\ a_2^{[1]} \end{bmatrix}, \quad A^{[2]} = \begin{bmatrix} a_2^{[2]} \\ a_2^{[2]} \end{bmatrix} = \begin{bmatrix} \sigma(z_2^{[1]}) \\ \sigma(z_2^{[2]}) \end{bmatrix}$$

Solution:

(i) $\frac{\partial f(x)}{\partial x_1^{[1]}} = w_{11}^{[3]} \frac{\partial a_1^{[3]}}{\partial x_1^{[1]}} = w_{11}^{[3]} \sigma(z_1^{[2]}) (1 - \sigma(z_1^{[2]})) = w_{11}^{[3]} a_1^{[2]} (1 - a_1^{[2]})$

$f = \begin{bmatrix} w_{11}^{[3]} & w_{12}^{[3]} \end{bmatrix} A^{[3]}, \quad A^{[2]} = \sigma(Z^{[2]})$

(ii) $\frac{\partial f}{\partial Z^{[2]}} = \frac{\partial f}{\partial A^{[2]}} \frac{\partial A^{[2]}}{\partial Z^{[2]}} = \begin{bmatrix} w_{11}^{[3]} & w_{12}^{[3]} \end{bmatrix}^T \circ A^{[2]} \circ (1 - A^{[2]})$

(iii) $\frac{\partial f(x)}{\partial Z^{[1]}} = \frac{\partial f(x)}{\partial Z^{[2]}} \frac{\partial Z^{[2]}}{\partial Z^{[1]}} = \frac{\partial f(x)}{\partial Z^{[2]}} \frac{\partial Z^{[2]}}{\partial A^{[1]}} \frac{\partial A^{[1]}}{\partial Z^{[1]}} = \begin{bmatrix} w_{11}^{[2]} & w_{12}^{[2]} \\ w_{21}^{[2]} & w_{22}^{[2]} \end{bmatrix} \delta_2 \circ A^{[1]} \circ (1 - A^{[1]})$

(iv) $\delta_1 = \frac{\partial f(x)}{\partial w_{11}} = \frac{\partial f(x)}{\partial Z^{[1]}} \frac{\partial Z^{[1]}}{\partial w_{11}} = \delta_1^T \begin{bmatrix} x_1 \\ 0 \end{bmatrix}$

GAN

Generative Adversarial Networks (GANs)

- GANs represent a game between two players:
 - Generator G: producing samples hard to distinguish from real training samples
 - Discriminator D: distinguishing G's generated samples from real training samples
- GAN cost functions:
 - Discriminator cost: $J^D = -\frac{1}{m_{\text{real}}} \sum_{i=1}^{m_{\text{real}}} \log(D(x^{(i)})) - \frac{1}{m_{\text{gen}}} \sum_{i=1}^{m_{\text{gen}}} \log(1 - D(G(z^{(i)})))$
 - Generator cost: a) Saturating cost $J^G = \frac{1}{m_{\text{gen}}} \sum_{i=1}^{m_{\text{gen}}} \log(1 - D(G(z^{(i)})))$
 b) Non-saturating cost $J^G = -\frac{1}{m_{\text{gen}}} \sum_{i=1}^{m_{\text{gen}}} \log(D(G(z^{(i)})))$

CNN

Convolutional Neural Networks (CNNs)

Convolution Layers:
Hyperparameters: Stride (s), Padding (p), Filter size (f)

- 2D (No depth)
- 3D (e.g. RGB channels)

learnable parameters = $(f^2 + 1 \text{ (for bias)}) \cdot n_i$
 Output shape = $(n-f+1) \cdot (n-f+1) \cdot n_i$

learnable parameters = $(f^2 + 1 \text{ (for bias)}) \cdot n_i$
 Output shape = $(n-f+1) \cdot (n-f+1) \cdot n_i$

Convolutional Neural Networks (CNNs)

Padding and Strided Convolution Layers:

- "Valid" convolutions \Leftrightarrow No padding:
 Output shape = $(n-f+1) \cdot (n-f+1) \cdot n_i$
- "Same" convolutions \Leftrightarrow Output shape should match input shape
 Output shape with padding = $(n-f+1+2p) \cdot (n-f+1+2p) \cdot n_i$
 p for "Same" convolutions = $(f-1)/2$
- General formula for output shape: $\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times n_i$

Pooling Layers:

Hyperparameters: Stride (s), Filter size (f)

- General formula for output shape: $\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times n_i$

Dropout

To do this, the authors suggest scaling the activation function by a factor of q during the test phase in order to use the expected output produced in the training phase as the single output required in the test phase. Thus:

Train phase: $O_i = X_i \alpha (\sum_{k=1}^d w_k x_k + b)$
Test phase: $O_i = q \alpha (\sum_{k=1}^d w_k x_k + b)$

A slightly different approach is to use **Inverted Dropout**. This approach consists in the scaling of the activations during the training phase, leaving the test phase untouched.

The scale factor is the inverse of the keep probability: $\frac{1}{1-p} = \frac{1}{q}$, thus:

Train phase: $O_i = \frac{1}{q} X_i \alpha (\sum_{k=1}^d w_k x_k + b)$
Test phase: $O_i = \alpha (\sum_{k=1}^d w_k x_k + b)$

Bias

Bias & Variance Recipe